

Visualization

Topics

1. Matplotlib core framework
2. Pandas plot()
3. Seaborn statistical visualization
4. (not covered) Grammar of graphics (ggplot2 see plotnine)
5. (not covered) Interactive plotting

Resources

1. Ch 9 in Python for Data Analysis, 2nd Ed, Wes McKinney (Ucalgary library and <https://github.com/wesm/pydata-book> (<https://github.com/wesm/pydata-book>))
2. Ch 4 in Python Data Science Handbook, Jake VanderPlas (Ucalgary library and <https://github.com/jakevdp/PythonDataScienceHandbook> (<https://github.com/jakevdp/PythonDataScienceHandbook>))
3. Fundamentals of Data Visualization, Claus O. Wilke (Ucalgary library and <https://serialmentor.com/dataviz/index.html> (<https://serialmentor.com/dataviz/index.html>))
4. Overview by Jake VanderPlas <https://www.youtube.com/watch?v=FytuB8nFHPQ> (<https://www.youtube.com/watch?v=FytuB8nFHPQ>)

Matplotlib

Matplotlib is a Python 2D plotting library which produces publication quality figures in a variety of hardcopy formats and interactive environments across platforms.

Matplotlib tries to make easy things easy and hard things possible.

For simple plotting the pyplot module provides a MATLAB-like interface

<https://matplotlib.org> (<https://matplotlib.org>)

Importing matplotlib looks like this

```
In [59]: %matplotlib inline

import numpy as np
import pandas as pd
import matplotlib as mpl
import matplotlib.pyplot as plt
```

Two interfaces

There are two ways to interact with Matplotlib: a Matlab style and an object oriented style interface.

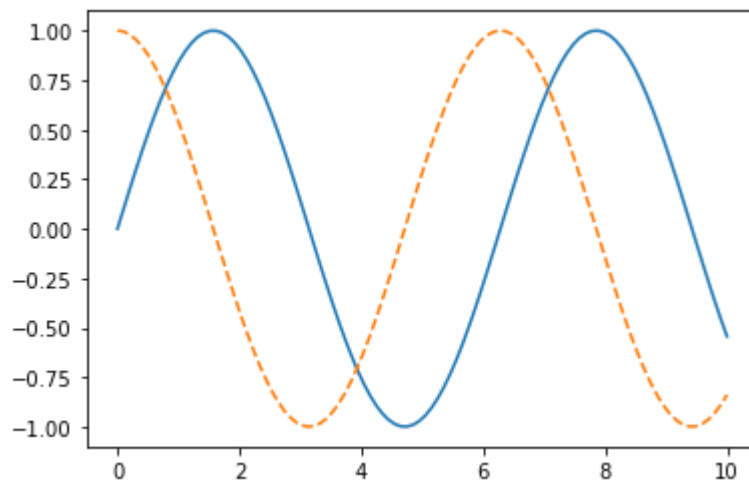
See Ch 4 in Python Data Science Handbook, Jake VanderPlas

- Two Interfaces for the Price of One, pp. 222
- Matplotlib Gotchas, pp. 232

Matlab style interface

```
In [60]: x = np.linspace(0, 10, 100)

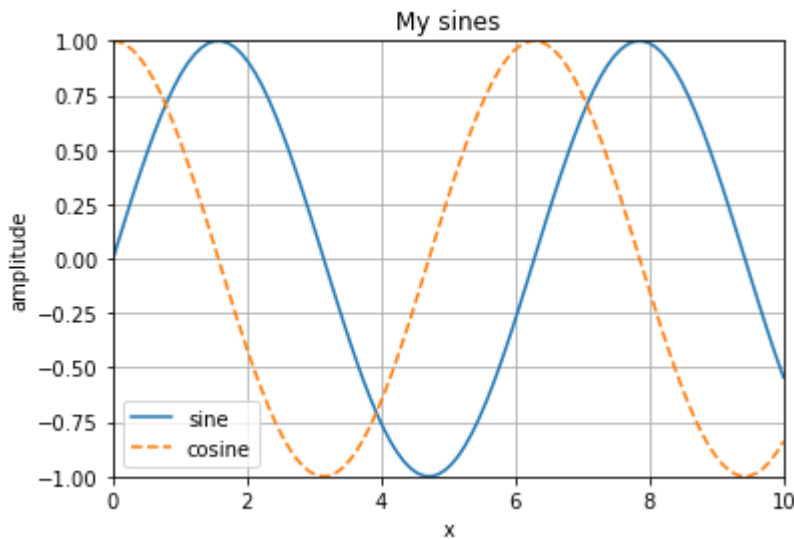
plt.plot(x, np.sin(x), '-')
plt.plot(x, np.cos(x), '--');
```



Adding decorations to the plot is done by repeatedly calling functions on the imported `plt` module. All calls within the cell will be applied to the current figure and axes.

```
In [61]: plt.plot(x, np.sin(x), '-', label='sine')
plt.plot(x, np.cos(x), '--', label='cosine')

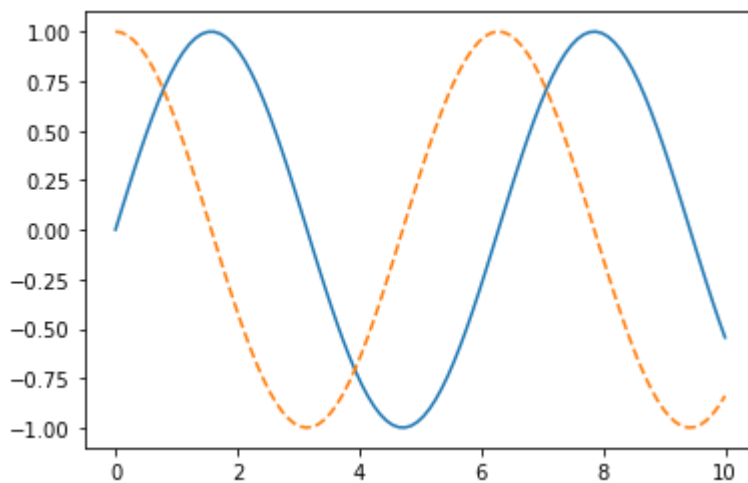
plt.xlim([0, 10])
plt.ylim([-1, 1])
plt.xlabel('x')
plt.ylabel('amplitude')
plt.title('My sines')
plt.grid()
plt.legend();
```



Object oriented interface

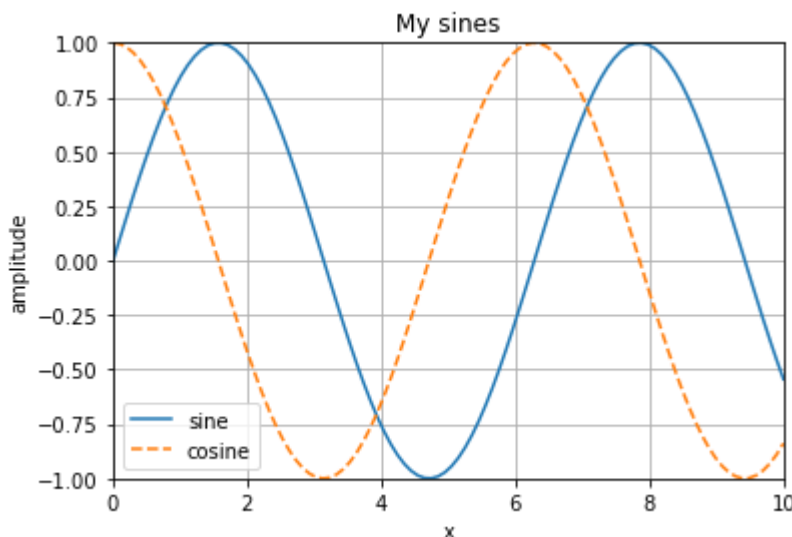
With this interface, you first create a figure and an axes object, then call their methods to change the plot.

```
In [62]: fig = plt.figure()
ax = plt.axes()
ax.plot(x, np.sin(x), '-')
ax.plot(x, np.cos(x), '--');
```



```
In [63]: fig = plt.figure()
ax = plt.axes()
ax.plot(x, np.sin(x), '-', label='sine')
ax.plot(x, np.cos(x), '--', label='cosine')

ax.set(xlim=[0, 10], ylim=[-1, 1],
       xlabel='x', ylabel='amplitude',
       title='My sines');
ax.grid()
ax.legend();
```



Save to file

With the figure object at hand, we can save to file

```
In [64]: fig.savefig('sines.pdf')
```

Plotting with pandas

We use the standard convention for referencing the matplotlib API ... We provide the basics in pandas to easily create decent looking plots.

https://pandas.pydata.org/pandas-docs/stable/user_guide/visualization.html
(https://pandas.pydata.org/pandas-docs/stable/user_guide/visualization.html)

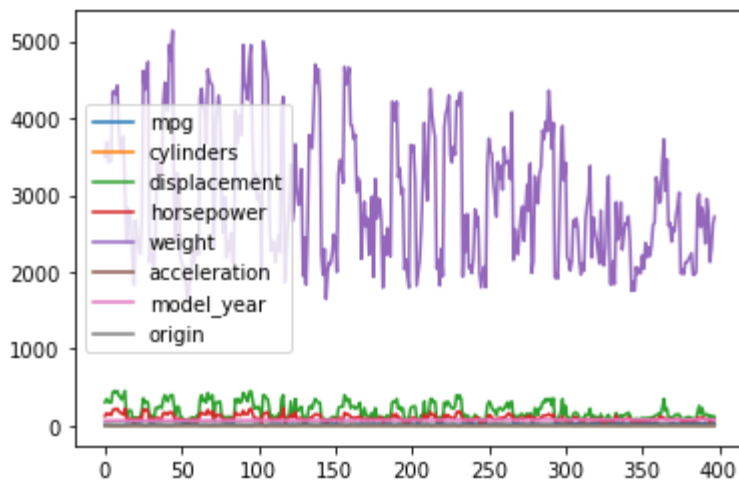
Let's load the heart attack dataset

```
In [66]: data = pd.read_csv('auto-mpg.csv', na_values='?')
```

Plotting all columns, works, but does not provide a lot of insight.

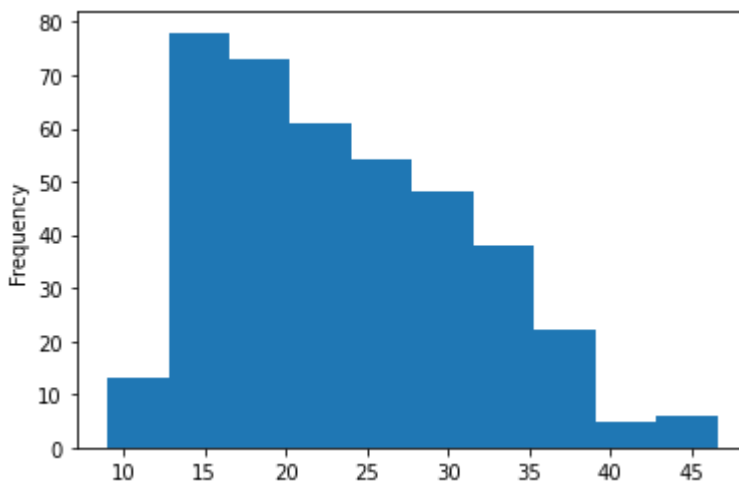
```
In [67]: data.plot()
```

```
Out[67]: <AxesSubplot:>
```



Let's look at the age distribution (a histogram)

```
In [68]: data['mpg'].plot.hist();
```



How many male and female samples do we have?

```
In [69]: data.origin.value_counts()
```

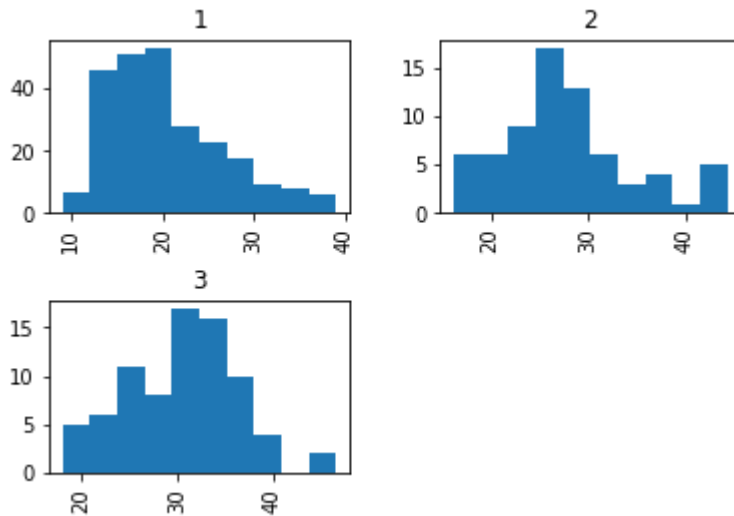
```
Out[69]: 1    249
         3     79
         2     70
         Name: origin, dtype: int64
```

Notice that we accessed the gender column with dot notation. This can be done whenever the column name is 'nice' enough to be a python variable name.

Do we have similar ages in females and males?

Plotting the histograms for each gender side-by-side directly from the dataframe:

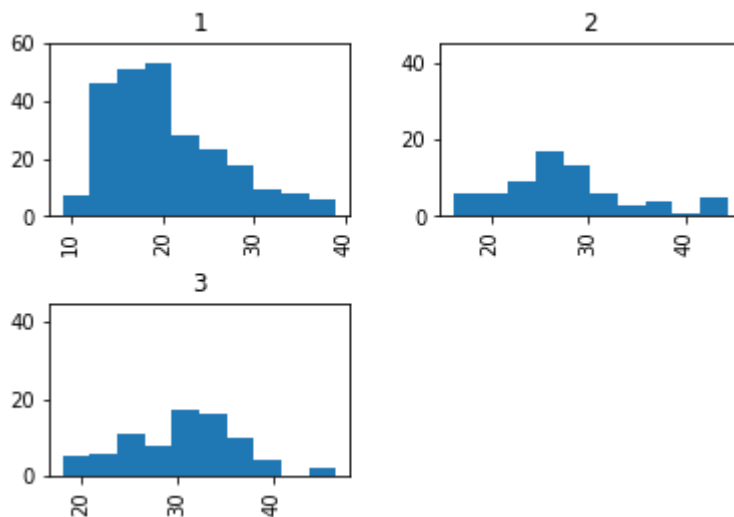
```
In [70]: axs = data.hist(column='mpg', by='origin')
```



To format this plot, we can work on the axes (array) that is returned by the plot call. We use Matplotlib object oriented interface methods to do this

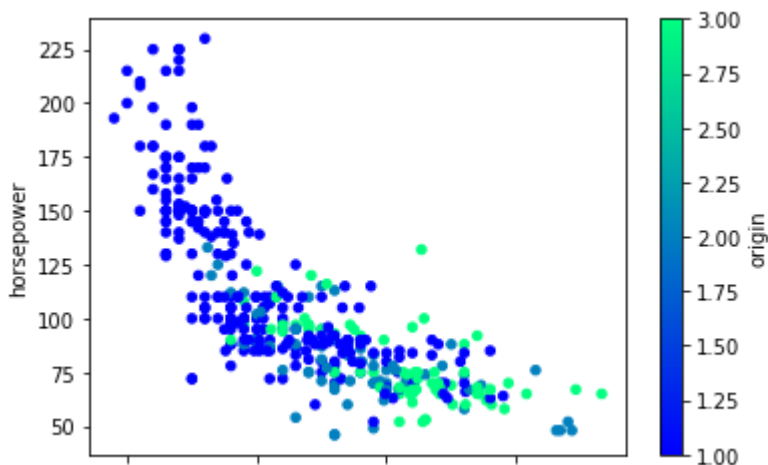
```
In [71]: axs = data.hist(column='mpg', by='origin')

axs[0,0].set(title='1', ylim=[0, 60])
axs[0,1].set(title='2', ylim=[0, 45]);
axs[1,0].set(title='3', ylim=[0, 45]);
```



Is age and blood pressure correlated? Maybe it is different for females and males?
Let's have a look with a scatter plot.

```
In [72]: data.plot.scatter('mpg', 'horsepower', c='origin', colormap='winter');
```



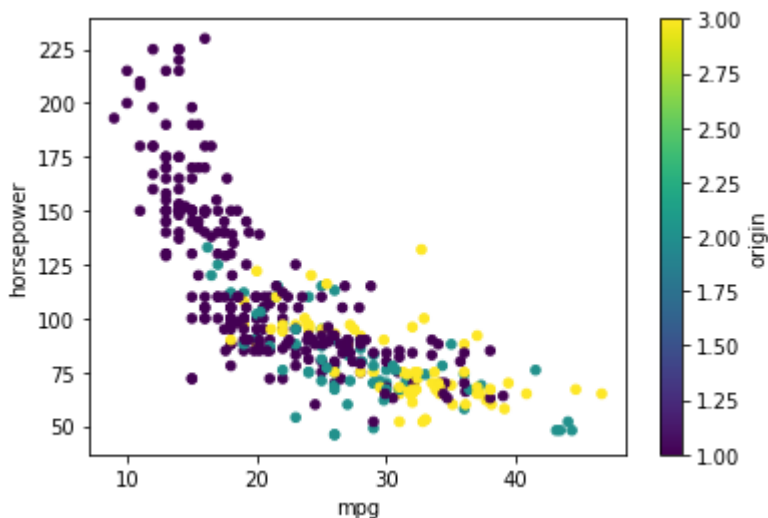
According to:

<https://stackoverflow.com/questions/43578976/pandas-missing-x-tick-labels>
(<https://stackoverflow.com/questions/43578976/pandas-missing-x-tick-labels>)

the missing x-labels are a pandas bug.

Workaround is to create axes prior to calling plot

```
In [73]: fig, ax = plt.subplots()
data.plot.scatter('mpg', 'horsepower', c='origin', colormap='viridis', ax=ax);
```

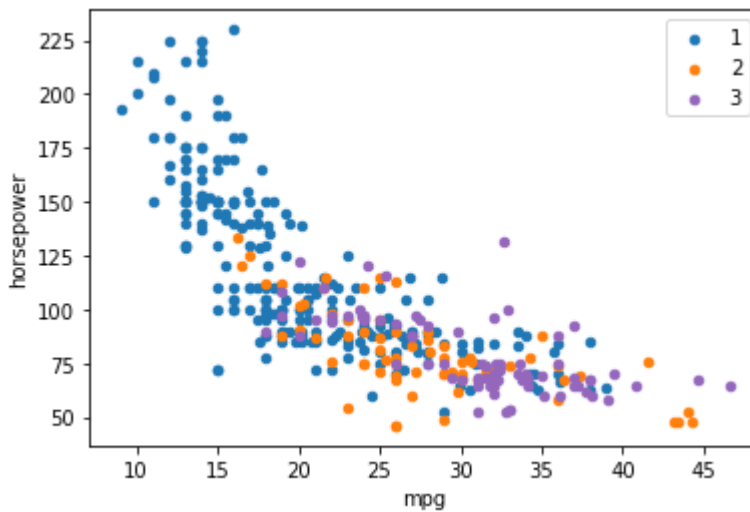


It is a bit annoying that there is a colorbar, we know gender is categorical.

One way to avoid the colorbar is to loop over the categories and assign colors based on the category.

See: <https://stackoverflow.com/questions/26139423/plot-different-color-for-different-categorical-levels-using-matplotlib> (<https://stackoverflow.com/questions/26139423/plot-different-color-for-different-categorical-levels-using-matplotlib>)

```
In [74]: colors = {1: 'tab:blue', 2: 'tab:orange', 3: 'tab:purple'}
fig, ax = plt.subplots()
for key, group in data.groupby(by='origin'):
    group.plot.scatter('mpg', 'horsepower', c=colors[key], label=key, ax=ax);
```



Seaborn

Seaborn is a Python data visualization library based on matplotlib. It provides a high-level interface for drawing attractive and informative statistical graphics.

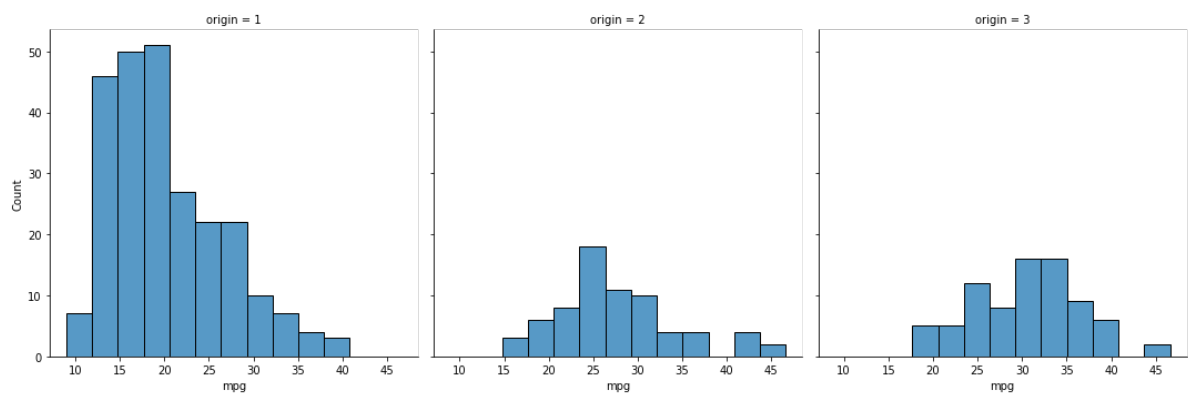
<http://seaborn.pydata.org/index.html> (<http://seaborn.pydata.org/index.html>)

Seaborn is usually imported as `sns`

```
In [75]: import seaborn as sns
```

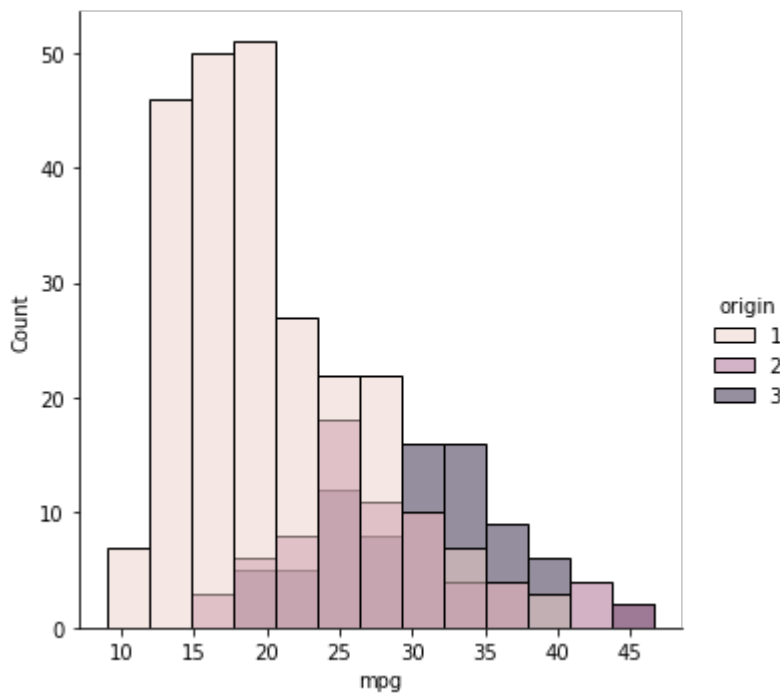
Let's re-create the histograms by gender with seaborn with the figure level `displot()` function.

```
In [76]: # Use gender to split age into columns
sns.displot(x='mpg', col='origin', data=data);
```



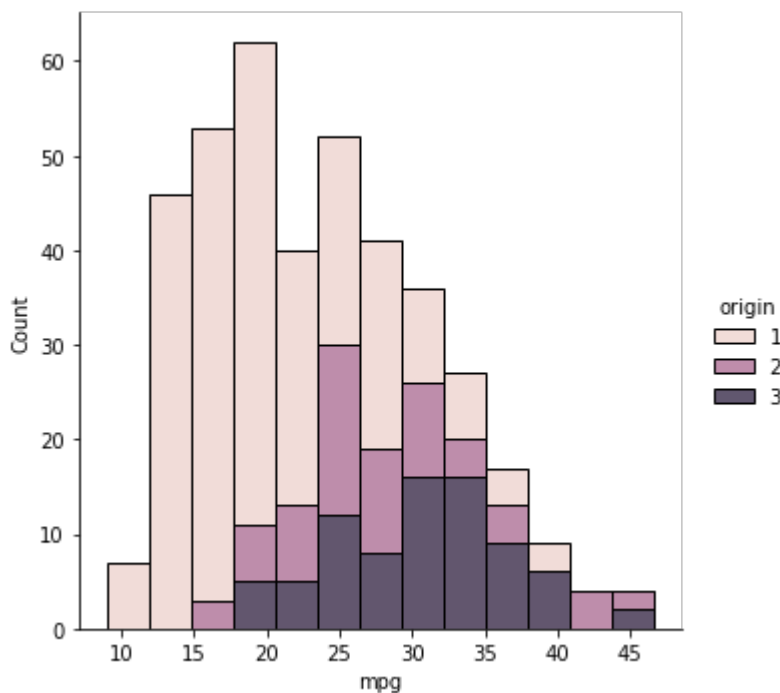
We can display the counts in the same plot, one on top of the other.

```
In [77]: # Use gender to color (hue) in the same plot
sns.displot(x='mpg', hue='origin', data=data);
```



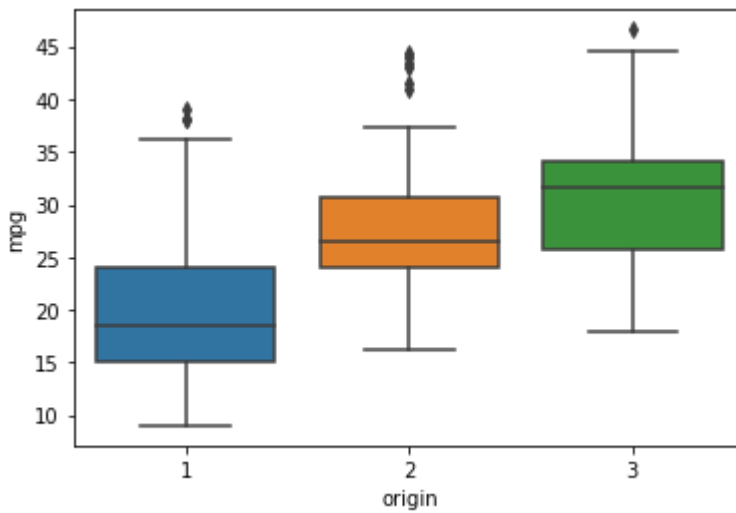
To have an idea of the split between male and female, we can stack the counts, adding up to total.

```
In [78]: sns.displot(x='mpg', hue='origin', data=data, multiple='stack');
```



We can look at the differences in ages with a boxplot too

```
In [79]: sns.boxplot(x='origin', y='mpg', data=data);
```

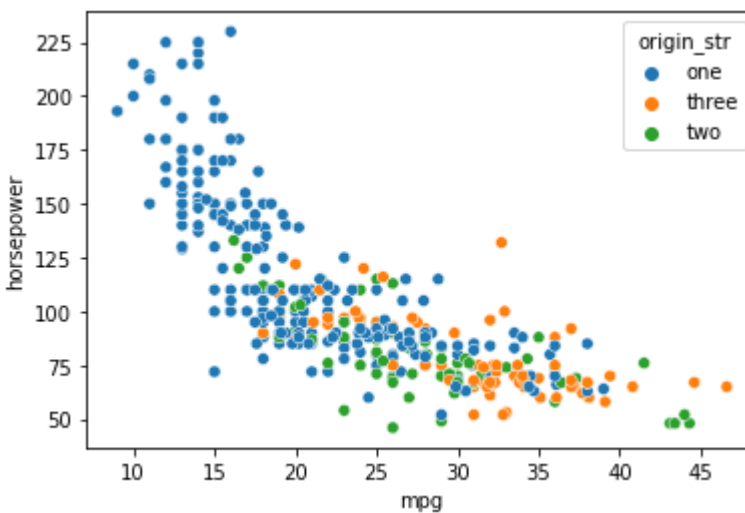


Let's re-create the scatter plot to see if age and blood pressure are correlated by gender.

To make the legend show strings we will create a gender string column with female and male strings rather than 0 and 1.

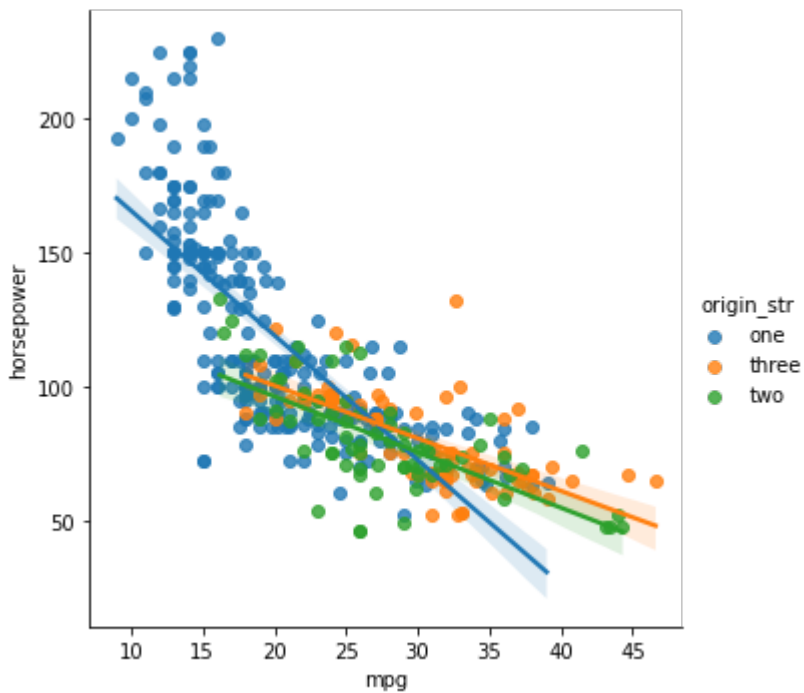
```
In [80]: data['origin_str'] = data['origin'].replace([1, 2, 3], ['one', 'two', 'three'])
```

```
In [81]: ax = sns.scatterplot(x='mpg', y='horsepower', data=data, hue='origin_str')
```



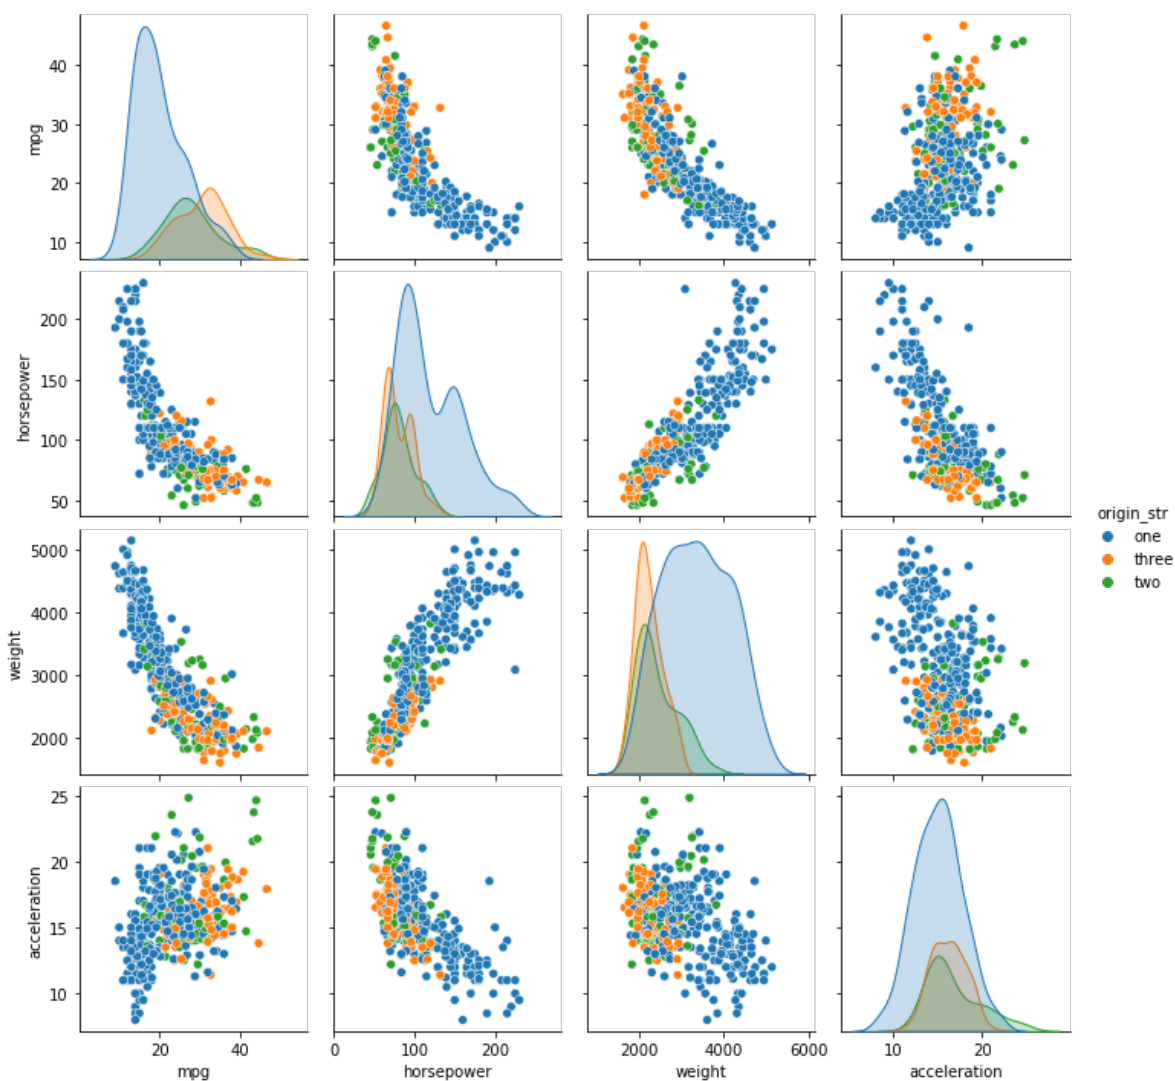
Adding a regression line helps with visualizing the relationship

```
In [82]: ax = sns.lmplot(x='mpg', y='horsepower', data=data, hue='origin_str')
```



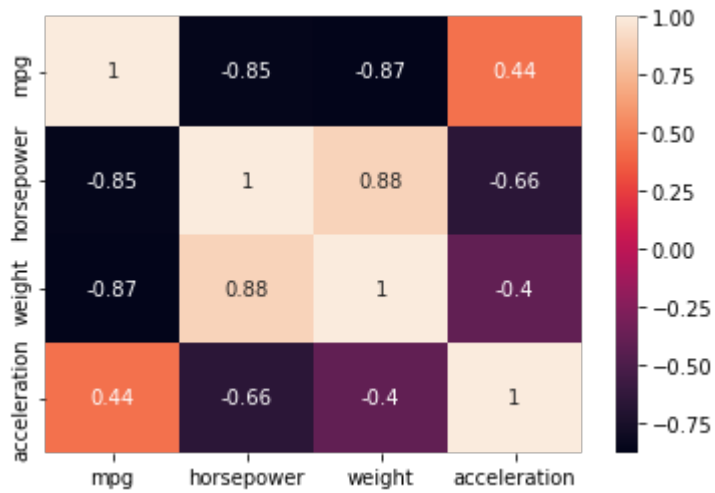
Maybe there are other correlations in the data set. Pairplot is a great way to get an overview

```
In [83]: sns.pairplot(data, vars=['mpg', 'horsepower', 'weight', 'acceleration'], hue='
```



As an alternative, we can visualize the correlation matrix as a heatmap

```
In [84]: g = sns.heatmap(data[['mpg', 'horsepower', 'weight', 'acceleration']].corr(method='spearmanr', annot=True))
```



There are nice tutorials on the Seaborn website, be sure to check these out.

```
In [ ]:
```